

# Real-Time Systems

[Galvin and others]

# Real-Time Systems

- System Characteristics
- Features of Real-Time Systems
- Implementing Real-Time Operating Systems
- Real-Time CPU Scheduling
- An Example: VxWorks 5.x

# Objectives

- To explain the timing requirements of real-time systems
- To distinguish between hard and soft real-time systems
- To discuss the defining characteristics of real-time systems
- To describe scheduling algorithms for hard real-time systems

# Overview of Real-Time Systems

- A **real-time system** requires that results be produced within a specified deadline period
- An **embedded system** is a computing device that is part of a larger system (I.e. automobile, airliner)
- A **safety-critical system** is a real-time system with catastrophic results in case of failure
- A hard real-time system guarantees that real-time tasks be completed within their required deadlines
- A **soft real-time system** provides priority of real-time tasks over non real-time tasks



# System Characteristics

- Single purpose
- Small size
- Inexpensively mass-produced
- Specific timing requirements

# System-on-a-Chip

- Many real-time systems are designed using system-on-a-chip (SOC) strategy
- SOC allows the CPU, memory, memory-management unit, and attached peripheral ports (I.e. USB) to be contained in a single integrated circuit

# Bus-Oriented System and System On a Chip (SOC)



# Features of Real-Time Kernels

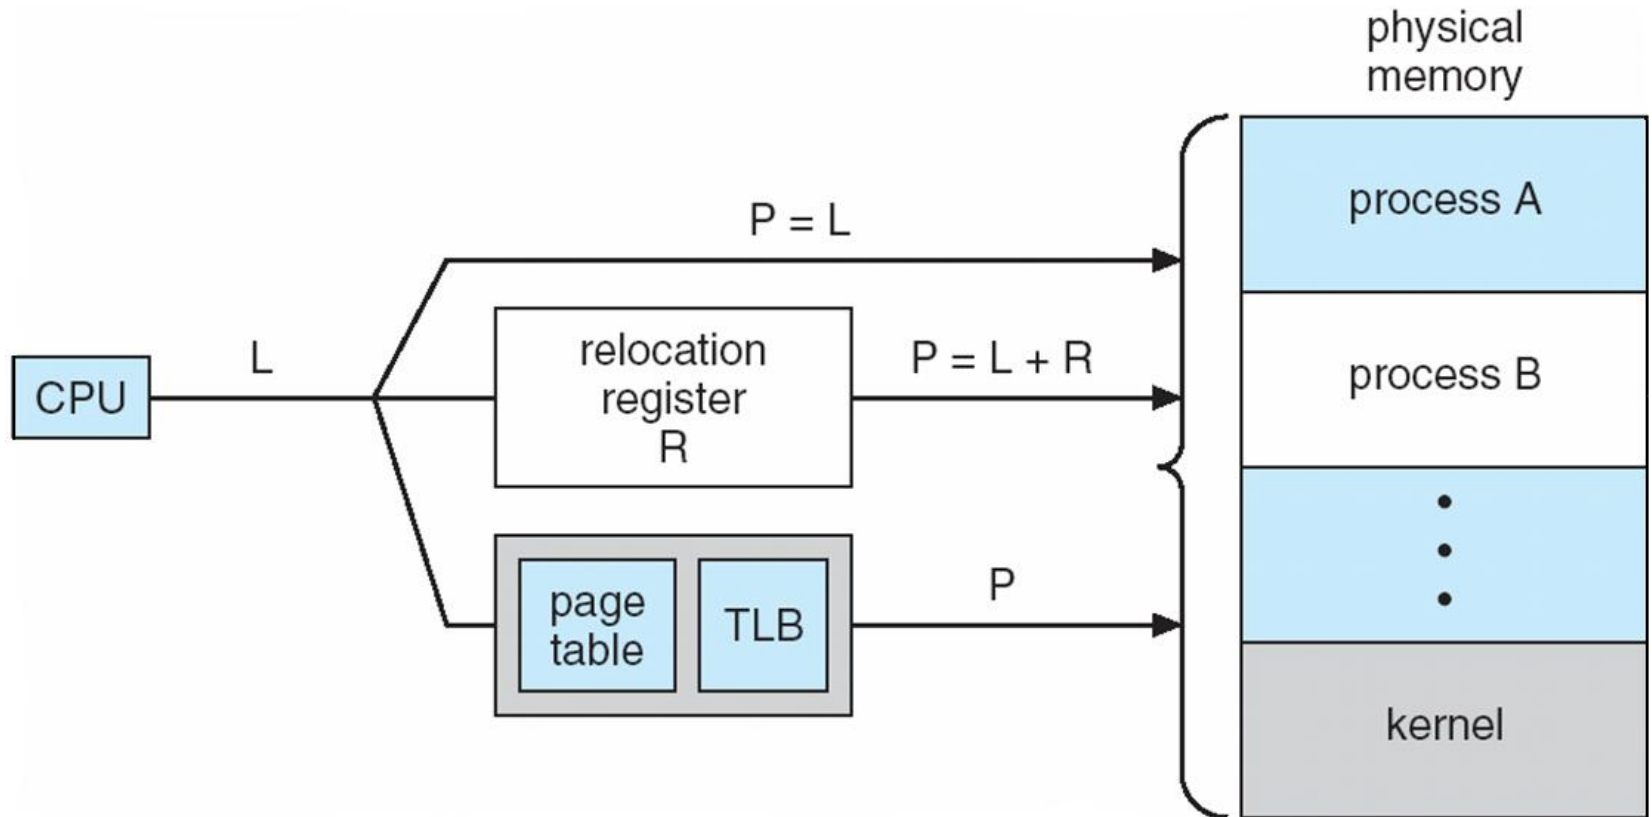
- Most real-time systems do not provide the features found in a standard desktop system
  - Multiuser interface
  - Peripheral Device Support
  - Protection and Security
- Reasons include
  - Real-time systems are typically single-purpose
  - Real-time systems often do not require interfacing with a user
  - Features found in a desktop PC require more substantial hardware than what is typically available in a real-time system



# Virtual Memory in Real-Time Systems

- Address translation may occur via:
  - (1) **Real-addressing mode** where programs generate actual addresses
  - (2) **Relocation** register mode
  - (3) Implementing full **virtual memory**

# Address Translation

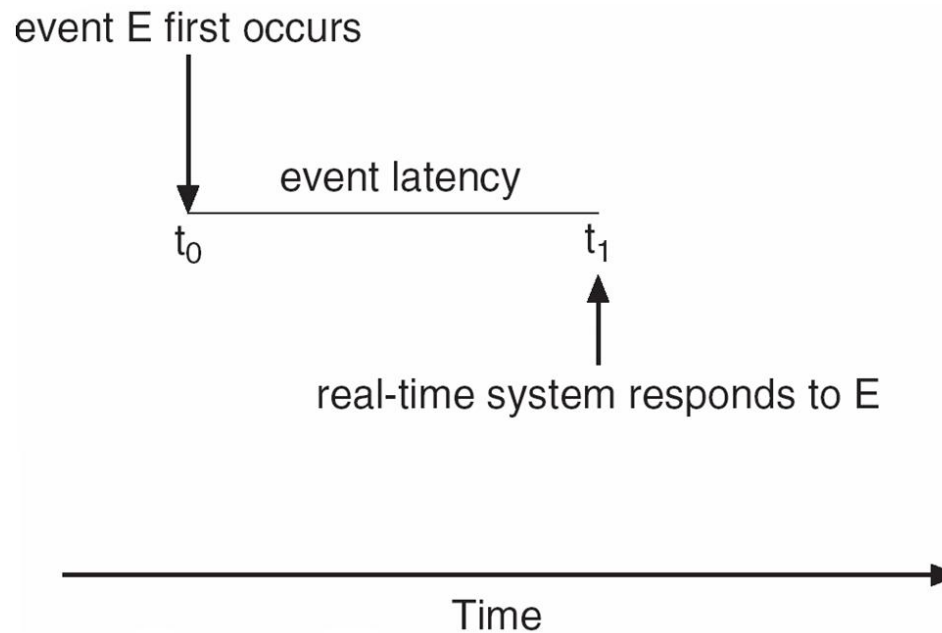


# Implementing Real-Time Systems

- In general, real-time operating systems must provide:
  - (1) Preemptive, priority-based scheduling
  - (2) Preemptive kernels
  - (3) Latency must be minimized

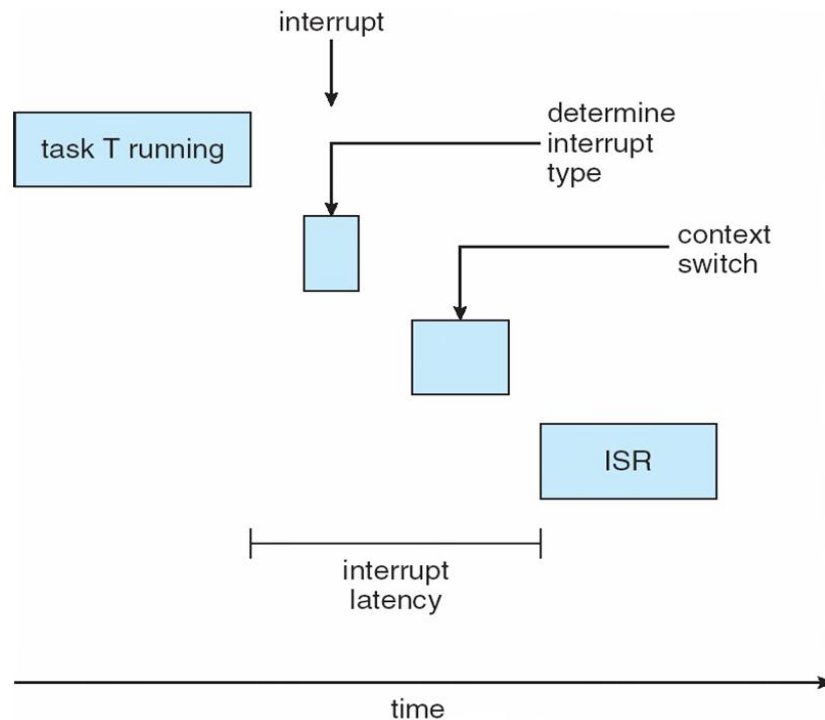
# Minimizing Latency

- **Event latency** is the amount of time from when an event occurs to when it is serviced.



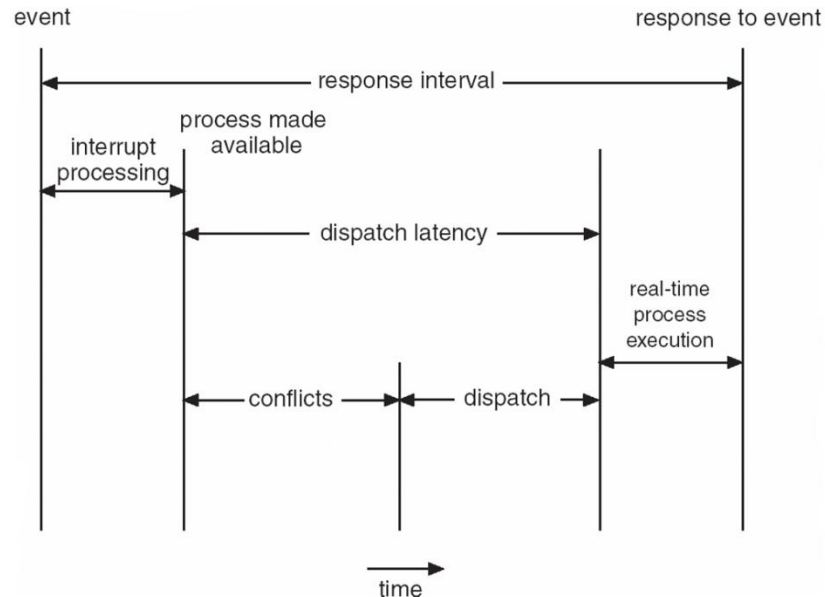
# Interrupt Latency

- Interrupt latency is the period of time from when an interrupt arrives at the CPU to when it is serviced




# Dispatch Latency

- **Dispatch latency** is the amount of time required for the scheduler to stop one process and start another




# Real-Time CPU Scheduling

- Periodic processes require the CPU at specified intervals (periods)
- $p$  is the duration of the period
- $d$  is the deadline by when the process must be serviced
- $t$  is the processing time



## Scheduling of tasks when $P_2$ has a higher priority than $P_1$

- $P_1 = 50$ ,  $t_1 = 20$
- $P_2 = 100$ ,  $t_2 = 35$







# Rate Monotonic Scheduling

- STATIC
- A priority is assigned based on the inverse of its period
- Shorter periods = higher priority;
- Longer periods = lower priority
- $P_1$  is assigned a higher priority than  $P_2$ .

- $P_1 = 50, t_1 = 20$
- $P_2 = 100, t_2 = 35$



- $P_1 = 50, t_1 = 25$
- $P_2 = 80, t_2 = 35$



Missed Deadlines with Rate Monotonic Scheduling


# Rate Monotonic Scheduling

- Worst Case Utilization is  $2(2^{1/n} - 1)$ ,  $n$ = no. of processes
- If RMS can not schedule a set of given processes then no other static priority scheduling algorithm can schedule them.

# Earliest Deadline First Scheduling

- Priorities are assigned DYNAMICALLY according to deadlines:
- the earlier the deadline, the higher the priority
- the later the deadline, the lower the priority

- $P_1 = 50, t_1 = 25$
- $P_2 = 80, t_2 = 35$



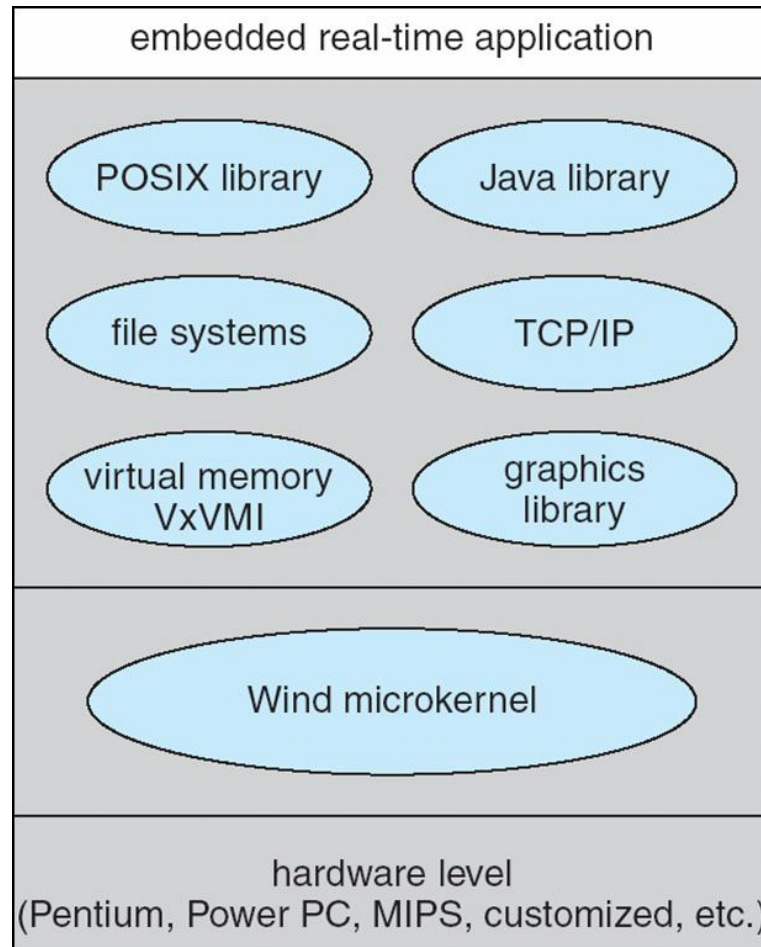
# Proportional Share Scheduling

- $T$  shares are allocated among all processes in the system
- An application receives  $N$  shares where  $N < T$
- This ensures each application will receive  $N / T$  of the total processor time

# Pthread Scheduling

- The Pthread API provides functions for managing real-time threads
- Pthreads defines two scheduling classes for real-time threads:
  - (1) `SCHED_FIFO` - threads are scheduled using a FCFS strategy with a FIFO queue. There is no time-slicing for threads of equal priority
  - (2) `SCHED_RR` - similar to `SCHED_FIFO` except time-slicing occurs for threads of equal priority

# VxWorks 5.0



# Wind Microkernel

- The Wind microkernel provides support for the following:
  - (1) Processes and threads
  - (2) preemptive and non-preemptive round-robin scheduling
  - (3) manages interrupts (with bounded interrupt and dispatch latency times)
  - (4) shared memory and message passing interprocess communication facilities

Thanks